



TikTok TechJam 2026 Information Document

*Note: This page is best viewed on desktop, and is accessible to the public via bit.ly/TikTokTechJam2026Info. Please feel free to share this with anyone interested in **TikTok TechJam 2026**.*



Be in the loop and get timely updates by joining our Telegram channel here:


[@TikTokTechJam2026](https://t.me/TikTokTechJam2026)



To successfully register for TikTok TechJam 2026, every individual must complete registration using **BOTH** links:

1. [Registration Form](https://bit.ly/TikTokTechJam2026Registration) <https://bit.ly/TikTokTechJam2026Registration>
2. [Devpost](https://tiktoktechjam2026.devpost.com/) <https://tiktoktechjam2026.devpost.com/>

The official deadline for registrations is **1 Sep 2026, 12pm**.



TikTok TechJam is back!


Now in its third year, TikTok TechJam 2026 is our flagship hackathon for university students across Singapore. Guided by our mission to *'Build with joy, code for change'*, this is your chance to break boundaries, unleash your creativity, and build solutions that drive real-world impact.

Build With Joy

'Build with joy' is a celebration of learning, collaboration, and curiosity. In line with TikTok's mission to *inspire creativity and bring joy*, participants are encouraged to embrace the spirit of innovation, where they can experiment freely, support each other, and have fun along the way as they grow together as builders of the future.

Code For Change

'Code for change' challenges participants to think beyond the code and focus on impact. We hope that through this hackathon, teams are empowered to build solutions that drive positive change, solve real-world problems, and reflect TikTok's belief in shaping the future with responsible technology.



This year, we're challenging you to go beyond basic prototypes. Across **5 distinct tracks**, you will tackle actual industry challenges spanning machine learning, e-commerce, and cybersecurity, while focusing on the robustness and efficiency required for real-world AI deployment.

But you won't be building alone. TechJam is as much about learning as it is about competing. Learn directly from our engineers through **virtual technical workshops**, and gain exclusive insights into the technologies that power our products.

Fly solo or form a team of up to 5 members to showcase your coding capabilities and compete for exciting prizes. Plus, you'll have the chance to connect directly with our hiring teams. This might just be your first step to launching your career at TikTok!

Who is this Hackathon for?

1. Students enrolled in a Singapore university
2. Students with an expected graduation date in December 2026 or later
3. Minimum age of 18 (as of the date of registration)
4. Solo participants, or teams of up to 5 members

Please find detailed rules & eligibility criteria on Devpost:

<https://tiktoktechjam2026.devpost.com/rules>.

Note that current interns and employees of TikTok are ineligible for participation in TikTok TechJam 2026.

Key Dates

- 23 August, 11:59pm: Early Bird Registration Deadline
- 25 August, 12pm: Problem Statements Released to Early Bird Only
- 27 August, 12pm: Problem Statements Released to Public
- 28 August: Tech Workshops (Webinars)
- 29 August, 12pm to 1 September, 12pm: 72-Hour Challenge & Project Submissions on Devpost
- 1 September, 12pm: Project Submissions Deadline on Devpost
- 8 September: Announcement of Finalists
- 11 September, 9am to 6pm: TikTok TechJam 2026 Grand Final @ TikTok Singapore
- 15 September: Winners Announcement on Devpost

How To Register

To successfully register for TikTok TechJam 2026, you must complete registration on **BOTH** our Registration Form & Devpost Page. Click the buttons below to register now!

Registration Form

Devpost

Virtual Technical Workshops (28 Aug, Friday, 1-6pm)

To set you up for success, we've lined up a **range of workshop webinars** led by our engineers — the very teams who created the problem statements for TikTok TechJam 2026! Look forward to key takeaways about the problems, our technologies, and the potential impact your projects can bring. You are highly encouraged to attend all the workshops that align with your interests, so as to make a more informed decision in selecting the track that works best for you.

- **Date:** 28 August 2026, Friday (SGT/GMT+8)
- **Time:** 1-6pm (see below for the actual 45min time slot of each track)
- **Webinar Link:** <https://vc-my.larkoffice.com/j/484622806>

Date	Time	Track
28 Aug (Fri)	1:00pm - 1:45pm	Track 1: Agent Launchpad: Design and Build Lightweight Agent Middleware
	2:00pm - 2:45pm	Track 2: Autonomous Machine Learning Research Agent for Recommender Systems
	3:00pm - 3:45pm	Track 3: Implement a GPU Kernel for a Transformer Layer
	4:00pm - 4:45pm	Track 4: Shopping Copilot: AI Conversational Search and Recommendations
	5:00pm - 5:45pm	Track 5: Robust Detection of AI-Generated Images Under Real-World Transformations



Stay updated by joining our Telegram channel here: [@TikTokTechJam2026](#)

Problem Statements



Join our **technical workshop webinars** to learn more about the tracks that interest you!

Look forward to key takeaways about the problems, our technologies, and the impact your projects can bring. You are highly encouraged to attend all the workshops that align with your interests, so as to make a more informed decision in selecting the track that works best for you.

- **Date:** 28 August 2026, Friday (SGT/GMT+8)
- **Time:** 1-6pm (see below for the actual 45min time slot of each track)
- **Webinar Link:** <https://vc-my.larkoffice.com/j/484622806>

1. Agent Launchpad: Design and Build Lightweight Agent Middleware



Technical Workshop Webinar with Q&A will be held on **28 Aug, 1:00 to 1:45pm**.

[Click here to join the webinar!](#)



Build the missing middleware, not the platform.

The **Starter Kit** already provides the browser UI, Agent CRUD, Playground, backend control plane, persistent workspaces, Codex CLI Runtime, BytePlus ModelArk integration, local containers, and optional ECS deployment. Identity and authorization, trace and audit, layered Agent architecture, and threat modeling and safety are recommended middleware examples—not a prescribed checklist. Teams may choose, combine, simplify, replace, or invent capabilities that make the Agent platform more usable, manageable, observable, secure, reliable, or extensible.

<https://github.com/RrankPyramid/CodeJam>

Starter Kit — RrankPyramid/CodeJam

1.1 Challenge Overview

AI Agents are software actors that can reason, call tools, execute code, read and write files, and continue work across multiple turns. A useful Agent platform therefore needs more than a chat box: operators must be able to understand what happened, control what an Agent may access, and contain unsafe execution.

Building the full web application, control plane, cloud deployment, model connection, and Agent Runtime from scratch would consume the entire hackathon. This challenge removes that bottleneck. Every team starts from the same working platform and spends the three days on one meaningful infrastructure problem.


Your goal: design and demonstrate a coherent middleware story that improves the Agent platform in a functional, testable way without breaking the provided lifecycle or Playground. Evaluation focuses on the relevance, quality, and integration of the capabilities your team chooses or invents.


Area	Provided by the Starter Kit	Student responsibility
Product experience	React UI, Agent list, Create/Edit forms, lifecycle controls, Playground, Run status.	Keep the baseline working; add only the UI needed to expose your middleware.
Control plane		

	Fastify API, validation, asynchronous Runs, AgentService, JSON persistence.	Integrate real middleware behavior into the backend path.
Agent Runtime	Codex CLI, persistent sessions, per-Agent workspaces, disposable local containers.	Integrate team-designed middleware at the most appropriate execution boundary.
Infrastructure	Docker, Colima, Podman, Docker Compose, ECS scripts, and Terraform.	Use the smallest runtime path that proves your design. Cloud deployment is optional.
Middleware	Intentionally absent: no user identity, trace timeline, audit model, or hardened sandbox policy.	Select, adapt, combine, or invent a coherent set of middleware capabilities and demonstrate why they improve the platform.

1.2 Starter Kit


What already works





- Create, inspect, edit, start, stop, and delete Agents from the browser.
- Send multi-turn tasks through the Playground and poll asynchronous Run status.
- Let Codex CLI write files and run commands inside the selected Agent workspace.
- Resume the same Codex session in later messages.
- Persist Agent, message, and Run metadata in a local JSON store.
- Run each local turn in a disposable Docker, Colima, or Podman container.
- Connect Codex to a BytePlus ModelArk Responses-compatible endpoint.
- Deploy the same POC to an existing BytePlus ECS instance or provision an ECS environment with Terraform.

Current architecture and extension points



The Fastify request boundary, `AgentService`, the `AgentRunner` interface, and the execution data model are all valid extension seams. The diagram deliberately uses one generic Team-Designed Middleware node: teams may add new events, principals, policies, lifecycle behavior, provider adapters, memory controls, reliability mechanisms, or other capabilities wherever their design has the strongest boundary.

Profile	Agent execution	Use during the hackathon
Local POC	One disposable local container per turn.	Recommended development and judging path. Supports Docker, Colima, and rootless Podman.
BytePlus ECS	Codex runs inside the application container.	Optional deployment path for teams that want a cloud demo.
Local development	Codex runs as a host process.	Useful for hot reload when the host Codex CLI is installed and configured.

Intentional limitations

The repository is a single-user POC. Its optional bearer token protects a remote demo but is not a user identity or authorization system. The JSON store supports one process. Ordinary containers are not a hardened multi-tenant isolation boundary. These limitations are deliberate extension points, not hidden requirements to fix all at once.

1.3 Run the Baseline Locally

Requirements

- macOS or Linux.
- Node.js 22 or newer and npm 10 or newer.
- One container engine: Docker, Colima, or Podman.
- A BytePlus ModelArk API key and a Responses-compatible endpoint ID.

Clone and start

Clone the Starter Kit

```
1  git clone https://github.com/RrankPyramid/CodeJam.git
2  cd CodeJam
```

One-command local POC

```
1 ARK_API_KEY=your-ark-api-key ARK_MODEL=ep-your-endpoint-id npm run poc
```

The first run installs Node.js dependencies and builds the Runtime image. The startup script automatically selects Docker, Colima, or Podman. Open <http://localhost:3000> when the server is ready.

Ark credential requirement: `ARK_API_KEY` must be an Ark model API key, not a BytePlus account AK/SK. `ARK_MODEL` is normally an endpoint ID beginning with `ep-`. A wrong credential produces a `401 Unauthorized` response from the Ark Responses API.

Force rootless Podman

```
1 CONTAINER_ENGINE=podman ARK_API_KEY=your-ark-api-key ARK_MODEL=ep-your-  
  endpoint-id npm run poc
```

Colima exposes the Docker CLI, so use the normal command after running `colima start`. Press `Ctrl+C` to stop the POC. Agent workspaces and conversations remain available for the next run.

Baseline acceptance test

1. Open the browser and select **Create Agent**.
2. Enter a name, description, and workspace instructions.
3. Create the Agent and send the following task in the Playground.

Example Playground task

```
1 Create a TypeScript hello-world CLI, add a test, run it, and summarize the  
  files you created.
```

1. Wait for the Run to complete and confirm that an assistant response appears.
2. Send a follow-up message and confirm that the same Codex session continues.
3. Stop and restart the Agent, then confirm that the workspace still exists.

Do not start middleware work until this flow succeeds. If the baseline fails, check `docker info` or `podman info`, inspect <http://localhost:3000/api/system>, and verify the Ark key and endpoint.

Development and validation

The simplest workflow is to edit the code, stop the POC, and rerun `npm run poc`. Before submitting, run the repository validation suite:

Required repository check

```
1 npm run check
```

This command runs TypeScript checks, server tests, and production builds. Additional setup paths are documented in the repository:

- [README and browser SOP](#)
- [Docker, Colima, and rootless Podman guide](#)
- [Architecture and extension seams](#)
- [Optional ECS deployment guide](#)

1.4 Platform and Middleware Design Requirements

The Starter Kit defines the basic platform experience. Beyond that baseline, middleware design is the core of the challenge. Teams should identify an Agent-specific problem, decide which responsibilities belong in the frontend, control plane, Runtime, data layer, or infrastructure boundary, and implement the smallest coherent solution that proves the idea.

The directions in Section 7 are examples rather than mandatory requirements. Teams may choose one, combine related ideas, simplify an example, replace it, or invent a different capability. Breadth is not the goal: reviewers will reward a clear problem, thoughtful architecture, real integration, and convincing evidence.

- **Preserve the baseline:** Agent CRUD, lifecycle actions, Playground chat, persistence, and model execution must continue to work.
- **Implement real behavior:** the middleware must execute in a backend, Runtime, data, or infrastructure path. Static screens and hard-coded success messages do not qualify.
- **Define the boundary:** explain which component owns the decision or event, what data crosses the boundary, and what happens when it fails.
- **Demonstrate meaningful evidence:** show the normal behavior and an appropriate failure, denial, recovery, degraded, or abuse case for the team's design.
- **Add automated verification:** test the core middleware behavior rather than only rendering the UI.
- **Keep secrets out:** never commit or display API keys, AK/SK, passwords, bearer tokens, or unredacted sensitive payloads.
- **Prefer the smallest useful infrastructure:** local execution is the default judging path; ECS is optional and does not affect the score.

In scope	Out of scope
A coherent middleware story with one or more related capabilities, a real integration path, minimal UI, tests, and demo evidence.	Rebuilding the React app, CRUD API, Playground, Codex integration, container launcher, or a commercial cloud platform.
Mock users, protected fixtures, controlled failures, provider adapters, lifecycle controls, trace data, policy decisions, or reliability mechanisms.	Production OAuth, a general-purpose policy engine, a microVM runtime, a container scheduler, or multi-region infrastructure unless central to the team's idea.
Focused schema changes and refactors needed to make the middleware understandable and extensible.	Unrelated redesigns or cosmetic work that does not prove Agent infrastructure behavior.

1.5 Agent Lifecycle and Post-Creation Experience

The Starter Kit already provides a post-creation experience rather than ending at an unexplained success message. A user can find an Agent, inspect its status and configuration, start or stop it, use the Playground, review messages and Runs, continue a Codex session, and delete the Agent while following an explicit workspace archival policy.

Teams may extend this lifecycle when it supports their middleware design. The following actions are examples rather than mandatory checkpoints:

- Test or invoke the Agent with a sample input.
- Open the middleware evidence for a specific Run, such as a trace, audit decision, policy result, recovery record, or budget event.
- Distinguish human operations from Agent operations and introduce human approval where useful.
- Update configuration through a new version and show what changed.
- Rotate or revoke credentials, permissions, tools, or network access.
- Pause, resume, stop, retry, reconcile, or recover an Agent or Run.
- Delete the Agent and clean up or retain its state according to an explicit policy.

Teams should implement only the lifecycle behavior needed to make their chosen or invented middleware capability convincing. Rebuilding every lifecycle feature is not expected.

1.6 Possible Three-Day Implementation Plan

Day	Engineering goal	Exit evidence
-----	------------------	---------------

1	Start and understand the baseline. Define the Agent-specific problem, choose or invent a coherent middleware story, specify the contract, and complete the first backend path.	The baseline passes; the team can trigger one real middleware behavior, event, decision, or control from an API or test.
2	Finish the core middleware path, persist its evidence, add the minimum UI, and implement the most important success and failure cases.	The complete scenario works end to end from the browser to the backend, Runtime, data, or infrastructure boundary.
3	Add automated tests, handle errors and cleanup, finish the architecture diagram and README, then rehearse the demo.	<code>npm run check</code> passes and the complete demonstration fits within three minutes.

1.7 Recommended Middleware Directions and Examples

Middleware design is a core part of this challenge. The examples below follow the same language and architecture boundaries as the original challenge brief. They are recommended examples—not a prescribed checklist. Teams may choose, combine, simplify, replace, or invent capabilities that better support their platform. Evaluation focuses on the relevance, quality, and integration of the capabilities the team designs.

Recommended Middleware Example: Identity and Authorization

Identity and authorization are one recommended middleware direction. Teams choosing this area may explore a distinction between a **human principal** and an **Agent principal**. For example, an Agent could use a separate identity instead of reusing a human session, personal access token, or shared platform credential.

Possible identity and authorization ideas include:

- **Human authentication:** identify the user who owns, creates, approves, updates, or stops an Agent.
- **Per-Agent identity:** create a distinct principal for an Agent or Agent version that can be rotated or revoked independently.
- **Delegated authority:** represent scoped, time-bound, and revocable permissions that describe what the Agent may access.
- **Policy enforcement:** perform authorization checks at a trusted backend, tool, data, or Runtime boundary rather than relying on UI restrictions.

- **Approval boundaries:** require optional human approval for selected external writes, high-cost actions, production operations, or sensitive data access.
- **Action attribution:** record the initiating human, executing Agent, requested scope, decision, target resource, and result.
- **Secret handling and revocation:** keep provider credentials on trusted backends, redact sensitive values, and demonstrate how later execution changes after revocation.

A small mock identity model is acceptable. For example, a team could show ownership isolation between User A and User B and prove that an Agent owned by User A cannot read User B's mock resource. A login screen without server-side authorization would not demonstrate the middleware itself.

Recommended Middleware Example: Trace, Audit, and Observability

Trace, audit, and observability are another recommended middleware direction. A team choosing this area could represent an Agent Run as a connected sequence of reasoning and actions rather than unrelated logs. Trace context may propagate across the frontend, control plane, Agent Runtime, model calls, tool calls, workspace operations, sandbox jobs, or cloud APIs that are relevant to the team's design.

Possible trace and audit ideas include:

- Stable identifiers such as Agent ID, Agent version, Run ID, session ID, trace ID, span ID, and actor type.
- Start time, duration, status, error details, and retry or cancellation relationships.
- Span categories such as orchestration, model call, tool call, memory access, sandbox execution, policy decision, human approval, or cloud operation.
- Inputs and outputs stored in a safely summarized or redacted form.
- Model, tool, Runtime, and infrastructure metadata needed to diagnose a Run.
- Token usage, cost, resource consumption, or other budget signals when available.

A trace-focused frontend could provide a Run list and a trace detail view with a tree or timeline, expandable spans, status filters, and a way to locate the failing step. A machine-readable query or export interface is an optional extension. Secrets and sensitive payloads should be redacted before storage or display.

Recommended Middleware Design Example: Layered Agent Architecture

Layered architecture is another recommended middleware design direction. Teams are encouraged to explain how responsibilities are organized, but no single layering model is required. The table below illustrates one possible architecture that can be simplified, adapted, or replaced.

Layer	Primary responsibility	Illustrative Starter Kit boundary
Experience Layer	Agent creation, catalog, Playground, middleware evidence, and lifecycle actions.	React Web UI calling stable platform APIs without holding the Ark key.
Control Plane	Agent specification, validation, status, Run orchestration, and reconciliation.	Fastify routes and <code>AgentService</code> .
Identity and Policy Plane	Human and Agent identity, delegation, approval, revocation, and audit.	A team-designed boundary around API, service, tool, or Runtime operations.
Agent Runtime Layer	Codex execution, model access, tool routing, retries, cancellation, and limits.	<code>AgentRunner</code> , local Runtime containers, or the ECS process.
Execution and Data Layer	Workspace files, persistent state, protected resources, connectors, and isolated execution.	Per-Agent workspaces, JSON metadata, mock services, or provider adapters.
Observability Layer	Trace ingestion, correlation, redaction, storage, query, visualization, and export.	New Run events, stores, APIs, and UI views added by the team.
Cloud Resource Layer	Compute, networking, storage, scheduling, and sandbox infrastructure.	Docker, Colima, Podman, or optional BytePlus ECS.

Teams may document the API or event contracts between selected layers and explain how their design could evolve to support another Runtime, identity provider, trace backend, tool, model, or infrastructure provider.

Recommended Middleware Example: Threat Modeling and Safety

Threat modeling and safety controls are another recommended middleware direction. Teams choosing this area could identify protected assets, actors, trust boundaries, abuse cases, implemented controls, and known residual risks. The table below provides examples; teams may focus on whichever threats are relevant to their design.

Threat	Controls to consider and demonstrate
Credential theft or exposure	Managed secret references, short-lived credentials, rotation, redaction, and exclusion of secrets from source, browser state, logs, and traces.

Privilege escalation or confused delegation	Least-privilege scopes, explicit delegation, backend policy checks, approvals, revocation, and complete actor attribution.
Prompt injection or tool misuse	Tool allowlists, typed schemas, target-resource scoping, output validation, execution limits, and approval for high-risk actions.
Sandbox escape or untrusted code	Non-privileged execution, restricted filesystems and networks, resource limits, controlled mounts, and patched Runtime images.
Cross-user access or data exfiltration	Ownership-aware authorization, storage isolation, scoped queries, outbound allowlists, protected metadata endpoints, and negative tests.
Runaway execution or cost	Timeouts, quotas, concurrency limits, maximum steps, token or cost budgets, and an administrative stop control.
Sensitive trace capture	Configurable capture levels, redaction before export, trace access control, and retention limits.

The Starter Kit's existing CPU, memory, PID, dropped-capability, and `no-new-privileges` defaults may be reused as baseline safeguards, but they do not by themselves constitute a new safety capability.

Recommended Middleware Example: Multi-Agent Coordination

Multi-Agent coordination is another recommended middleware direction. A team choosing this area may connect several Agent instances through a shared session, topic, queue, or lightweight coordinator. The purpose is not to build a complex distributed system; it is to demonstrate that the platform can route messages, preserve shared state, and coordinate turns across more than one Agent Runtime.

Example demo: create several Agents and ask them to count down from 10 to 1 in a shared conversation. On each turn, one Agent publishes the next unused number, then another Agent continues until the sequence reaches 1.


The minimum coordination layer may provide:

- A shared session or topic that all participating Agents can read and write.
- A simple turn-selection or message-routing rule.
- Shared state that records the latest number and prevents duplicate or skipped turns.
- A visible event history showing which Agent produced each number.
- A timeout, retry, or stop rule for an Agent that does not respond.

Successful demonstration: the team starts multiple Agents from the platform, launches one shared task, and shows a complete 10-to-1 sequence with no duplicate or missing number. The

interface should make the participating Agent and ordering of each message clear. A platform-local webpage is sufficient; integration with an external chat product is optional.

Illustrative reference: multiple Agents take turns counting down in one shared topic.



Other Team-Designed Middleware

Teams are encouraged to propose capabilities outside the examples when they improve the Agent platform in a clear way. Possible directions include lifecycle reconciliation and failure recovery, state and memory governance, human-in-the-loop workflows, cost and budget control, provider abstraction, versioning and rollback, multi-Agent coordination, tool or model routing, credential exchange, or automated diagnosis and remediation.

A team-defined capability should still explain the Agent-specific problem, architecture boundary, functional evidence, failure or recovery case, and known limitations.

1.8 Required Live Demo

The live demo should show one complete scenario. The essential journey is a user creating or selecting a runnable Agent from the frontend and then using or testing it. Beyond that baseline,

each team should demonstrate the middleware capabilities it chose or designed. Identity, trace, audit, policy, recovery, and revocation are examples rather than mandatory checkpoints.

1. Create or select an Agent from the frontend and show its current lifecycle state.
2. Invoke the Agent through the Playground with a real task.
3. Show at least one real model, file, tool, sandbox, data, or infrastructure action.
4. Demonstrate the middleware behavior and the evidence it produces.
5. Demonstrate an appropriate failure, denial, degraded, abuse, or recovery case.
6. Show that the platform remains understandable and controllable afterward.

The scenario may use a mock third-party service or controlled fixture. The frontend-to-Agent path and any middleware presented by the team must be functional rather than represented only by static screens.

1.9 Deliverables

1. **Three-minute live demo:** show one real Agent Run and the team-designed middleware working in its normal case and an appropriate failure, denial, recovery, degraded, or abuse case.
2. **One-page architecture diagram:** show the middleware, data flow, trust boundary, and enforcement, instrumentation, or recovery point.
3. **Code repository:** include setup instructions, the middleware problem and rationale, design summary, automated tests, demo steps, limitations, and no secrets.

1.10 Core Acceptance Checklist and Optional Evidence

- ☐ A reviewer can clone the repository, start the platform, and create or test an Agent from the frontend.
- ☐ The submission identifies and demonstrates one or more meaningful middleware capabilities selected, adapted, combined, or designed by the team.
- ☐ The middleware executes in a backend, Runtime, data, or infrastructure path rather than only in the UI.
- ☐ The repository and documentation are sufficient for reviewers to understand and reproduce the POC.
- ☐ `npm run check` passes.
- ☐ No secret appears in source, Git history, logs, traces, screenshots, browser storage, or demo output.
- ☐ Optional evidence: a delegated permission is scoped or revocable, enforced outside the UI, and demonstrated.

- ☐ Optional evidence: an end-to-end Agent Run produces a correlated trace with relevant model, tool, sandbox, policy, or infrastructure events.
- ☐ Optional evidence: a defined threat is blocked or contained, the protected asset remains unchanged, and cleanup or recovery is demonstrated.
- ☐ Optional evidence: a team-defined lifecycle, reliability, memory, budget, provider, or coordination capability works as described.

1.11 Evaluation Criteria

This track will follow the evaluation criteria below:

Category	Weight	What reviewers will look for
End-to-end middleware behavior	40%	A real frontend-to-backend, Runtime, data, or infrastructure path with convincing functional evidence.
Technical design and integration	25%	A clear rationale, coherent architecture, appropriate boundary, focused changes, and extensible contracts.
Verification and robustness	20%	Automated tests, error handling, cleanup or recovery, redaction, and protection against obvious bypasses.
Demo and reproducibility	15%	A concise live demo, useful README, one-command startup, documented limitations, and no hidden manual setup.

1.12 Scope Guidance and Frequently Asked Questions

This is a hackathon-scale Agent infrastructure challenge, not a requirement to build a complete commercial cloud product. A strong submission may support one local Runtime path, a small mock resource set, and a focused middleware story. Depth, coherence, and relevance matter more than the number of example features implemented.

Teams are not required to train a foundation model, build a workflow editor, implement production OAuth, create a general-purpose sandbox, support multiple cloud regions, or deploy to ECS. Mock external services are acceptable, but static UI mockups cannot replace functional middleware behavior.

Do we need BytePlus ECS? No. Local Docker, Colima, or Podman is the default judging path. Cloud deployment is optional.

Do we have to select one recommended example? No. The examples are starting points. Teams may adapt, combine, simplify, replace, or invent capabilities that fit their platform.

Can we use mock users or resources? Yes. Controlled fixtures are encouraged when they make middleware behavior reproducible.

Does a polished UI count as middleware? No. The UI may explain and visualize a capability, but the behavior must execute in a trusted backend, Runtime, data, or infrastructure path.

Why does Ark return 401 Unauthorized? The most common cause is using a BytePlus account AK/SK instead of an Ark model API key, or using the wrong endpoint ID.

Where should we start reading the code? Begin with `apps/server/src/types.ts`, `apps/server/src/app.ts`, `apps/server/src/agent-service.ts`, and the two `AgentRunner` implementations. Then inspect `apps/web/src/App.tsx` for the smallest UI integration point.

How to access the Starter Kit? <https://github.com/RrankPyramid/CodeJam>

2. Autonomous Machine Learning Research Agent for Recommender Systems

 In response to some queries from our Early Bird participants, our engineers have provided updates to the problem statement to improve clarity and to support participants better.

Problem Statement last updated: 27 August 2026, 5:55PM.

Added downloadable kuairand-starter-kit.zip under 'Starter Kit'

Problem Statement in our Early Bird release doc is also the same version as is here.

 **Technical Workshop Webinar with Q&A** will be held on **28 Aug, 2:00 to 2:45pm.**

[Click here to join the webinar!](#)

2.1 Background

Motivation

Machine learning engineers (MLEs) spend much of their time on a single activity: **taking a dataset and a set of metrics, then iterating on a model again and again to push the score higher.** This work is inherently cyclic — every round repeats the same loop, shown in Figure 1.




Figure 1. The MLE iteration loop

Figure 1. The MLE iteration loop. A closed cycle of five core stages, plus a reflection step that feeds the next round:

1. **Read the problem** — understand the given dataset and the target metrics.
2. **Inspect data** — study data distribution through exploratory data analysis (EDA).
3. **Engineer features** — build and select input features (see Appendix A.5).
4. **Train + tune** — choose a model, set the loss function, and tune hyperparameters.
5. **Evaluate** — read the metrics, check for overfitting, and consult the leaderboard.

The result of the **evaluate** stage drives a **reflect + revise** step, which decides what to change and loops back into the next iteration — re-inspecting the data and adjusting the features. The cycle repeats until the score plateaus.

Two of these stages — **engineer features** and **train + tune** — are carried out almost entirely in code: the engineer writes scripts to transform the data, define the model, and run training. In other words, each turn of the loop produces and modifies code. This is what makes the loop a natural target for automation: it is structured and repeatable, yet writing and revising that code is exactly the kind of task a code-generating LLM can take on.

The loop is also repetitive and mechanical. It draws heavily on "engineering intuition," but many individual steps are well-structured and repeatedly exercised in practice — which is precisely why automating the whole cycle has become an active research direction.

Prior Work

Over the past two years, a new line of work has set out to automate this loop: the **Autonomous ML Research Agent**, an LLM-driven agent that runs the cycle in Figure 1 on its own. It reads the

problem, **writes the code** for each stage, trains and evaluates the model, reflects on the results, revises its approach, and finally produces a submission. Representative systems include:

- **MLE-Bench** [1] (OpenAI) — a benchmark of 75 Kaggle competitions, now a standard evaluation suite for such agents.
- **AIDE** [2] (Weco AI) — a state-of-the-art agent that frames ML engineering as code optimization and explores the space of solutions via tree search.
- **AI-Scientist-v2** [3] (Sakana AI) — an end-to-end agent for autonomous scientific and ML research, using agentic tree search to form hypotheses, run experiments, and write up results.

This Challenge

This challenge asks participants to design an **autonomous ML research agent**. Given a public ML dataset and a set of metrics, the agent must **autonomously** run the full loop of Figure 1 — read the problem, engineer features, train and tune the model, evaluate, then reflect and iterate — to reach the highest possible score across the test sets. Writing the code for each stage is part of the agent's job, not something provided in advance.

New to recommender systems? All benchmarks in this challenge come from the recommendation domain (the KuaiRand family). If terms such as CTR, multi-task learning, GAUC, or NDCG are unfamiliar, start with the **Appendix: A Primer on Recommender Systems**. At the end of this document — a concept map plus an annotated reading list designed to get you oriented in 1–2 hours.

2.2 Problem Statement

The Task

Design and implement an Autonomous ML Research Agent. For each benchmark, the agent must autonomously:

1. **Reproduce the official baseline.** Stand up a working end-to-end pipeline and confirm it reaches the official baseline's reported validation score. (The official baseline is a fixed, organizer-provided reference — see *Benchmarks*. Any starter pipeline the agent builds for itself is an internal step, not the reference it is scored against.)
2. **Iterate on the pipeline.** Autonomously draw on established methods from both industry and academia to improve each stage of the pipeline (see Figure 1), and apply those improvements in code. The agent develops using **only the training split and the public validation feedback** — it never has access to the hidden test set.

3. **Improve over the baseline.** Through repeated iterations, drive the **validation** score above the official baseline. Improvement need not be strictly monotonic — as with real-world data, the trajectory may fluctuate — but the agent should show a clear, sustained ability to keep improving relative to the baseline. Final ranking is computed once, on the **hidden test set**, using the submission the agent designates as final.

Task Requirements

4. **Runs end-to-end and aims to beat the baseline.** The agent must run the full pipeline on the required benchmark (KuaiRand-Pure) and reach a converged result; attempting the bonus benchmark (KuaiRand-1k & KuaiRand-27k) is optional. The target is a hidden-test score that exceeds the official baseline; the actual delta achieved — positive or negative — is what feeds into the Primary metric scoring (see Judging Criteria), so falling short of the baseline is scored continuously rather than treated as a disqualifying failure.
5. **Iterates autonomously across the full stack.** The agent should improve the solution on its own, driven by its own evaluation of results. Improvements may target any part of the algorithmic stack — not just the model architecture, but every upstream and downstream module is fair game. The goal is to **minimize human intervention** — a fully autonomous run is the ideal, but a well-instrumented **semi-automated** pipeline that requires only a handful of interventions is an acceptable and realistic outcome; in practice, we measure how little human intervention a run requires (e.g. the number of manual interventions).
6. **Robust operation.** The pipeline should run reliably with **minimal human intervention**. Robustness here is about how the agent *handles* difficulty, not how often it succeeds — we do not score it by failure count, since a capable agent may fail only on genuinely hard problems. What matters is that when a step fails (a code error, a timeout, an unexpected input), the agent can recover, retry, or route around it, and that long iterative runs neither crash, stall, nor diverge.


2.3 Constraints & Scope

Category	Constraints & Scope Details
In scope	- Any open-source library or framework (PyTorch, RecBole, TorchRec, LightGBM, ...) - Any papers, public solutions, or pretrained weights - Changes to any pipeline stage — not just the model
Out of scope	- No external training data or pretrained weights trained on these benchmarks' test labels - No hidden-test access during development (train + validation only)

Limits	• KuaiRand-Pure: NDCG@10 / Recall@50, click = positive (fixed) (<i>Required</i>); KuaiRand-1k & KuaiRand-27k: same task and metrics (<i>Bonus</i>) • Hidden test scored once, on the final submission • Compute budget: 50 iterations per benchmark run (hard cap; the convergence rule $\epsilon = 0.002 / N = 3$ normally triggers first), plus a 6 h wall-clock ceiling per run as a backstop. Compute is deliberately not the binding constraint on this benchmark — 100 iterations of the official baseline take about 28 min on a single CPU core with no GPU. GPU-hours and LLM tokens are reported for Feasibility scoring, not capped.
Allowed assumptions	• Fixed <code>train / validation / hidden-test</code> split per dataset • Official baseline, scores & evaluation script (incl. convergence rule) • Example submission + output schema

2.4 Available Resources & Data

Starter Kit



To lower the barrier to entry — especially for participants new to recommender systems — the challenge provides a standard starting point. **Download: kuairand-starter-kit.zip** (above) — numpy only (no torch / pandas / scikit-learn); `python3 baseline.py --model fm` reproduces the official baseline in about 40 s on a single CPU core. It contains:

1. **Fixed data splits:** date-based, taken from the two standard logs
 (`log_standard_4_08_to_4_21_pure.csv` & `log_standard_4_22_to_5_08_pure.csv`). **train** = date 20220408–20220421 (1,141,112 rows) / **validation** = date 20220422–20220428 (124,909 rows) / **test** = date 20220429–20220508 (170,588 rows). Teams develop on train + validation only; the hidden test set is scored once. Splitting by date rather than by row count avoids any tie-breaking ambiguity on equal timestamps.
2. **Official baseline:** a fixed, organizer-provided reference pipeline shipped in the Starter Kit — a Factorization Machine ($k=16$, $lr=0.001$, 5 categorical fields), numpy only, about 40 s on CPU. Published **hidden-test** scores: GAUC **0.6610** / nDCG@5 **0.5282** / primary **0.5946** (mean over 5 seeds, std 0.0008). Validation: GAUC 0.6674 / nDCG@5 0.5357 / primary 0.6016. Reference rungs for harness self-check — random scoring: primary 0.4753; item popularity: primary 0.5715. Beating *this* baseline is what counts — not a baseline the team builds itself.

3. **Evaluation script:** the exact scoring code (GAUC / nDCG@5) ships in the Starter Kit as `evaluate.py`. It is model-agnostic — it takes only `(user_ids, labels, scores)`, so any model can be scored with it. **Pinned conventions:** users with zero positives count as $\text{nDCG} = 0$ and are included in the average; GAUC counts only users with $0 < \text{positives} < \text{impressions}$, weighted by positive count; $\text{nDCG gain} = 2^{\text{rel}} - 1$. **Convergence rule:** $\epsilon = 0.002$, $N = 3$ — a run is converged when the validation primary score has not improved by more than ϵ over the last N consecutive iterations ($\epsilon \approx 2.5\sigma$ of the baseline's 5-seed std of 0.0008). The absolute-delta aggregation is unchanged.
4. **Submission format:** a CSV with the header `row_id,user_id,video_id,score`, one line per evaluation-split row. `row_id` is a 0-based, strictly increasing index into the split as produced by `data.load()`; `user_id` / `video_id` are redundant fields used only to verify alignment; `score` is any real number (only the relative order matters), and NaN / Inf are rejected. The `row_id` is required because `(user_id, video_id)` is **not unique** in the evaluation split — 3.06% of test rows are repeated pairs, up to 12 times — so it cannot serve as a key. Generate a runnable example with `python3 submit.py --make` and validate with `--check`, which rejects a wrong header, a row-count mismatch, `row_id` gaps, misalignment against the evaluation split, and non-numeric scores.
5. **Run-log requirements:** each iteration should record its **hypothesis**, the **code diff**, the resulting **metrics**, and any **error / recovery events**. These logs are how judges assess **Autonomy** (scored under Impact & Relevance) and **Robustness** (scored under Technical Execution) — see Judging Criteria.
6. **LLM coding agent:** you can use whatever you like, or use [Trae](#) from ByteDance, which provides "Limited offer: new user 7-day free trial".

Benchmarks

KuaiRand-Pure is required and determines 100% of the primary score. **KuaiRand-1k and KuaiRand-27k are bonus datasets** — attempting them is optional and earns extra credit, but neither is required to complete the primary score.

Resource policy. This is a hackathon, so external resources are open by default: use any open-source library (PyTorch, RecBole, TorchRec, LightGBM, ...), read any papers, docs, or public solutions, and use pretrained model weights freely. The agent is expected to draw on whatever published methods it can find — that is what makes it a *research* agent.

There is **one hard rule: no external training data**. Training must rely only on the KuaiRand datasets listed below — no augmenting, joining, or pre-training on any other dataset, and no pretrained model whose weights were trained on these benchmarks' test labels. This single rule is what keeps the hidden-test ranking fair; everything else is unrestricted.

Dataset	Domain & Description	Metrics	Scale
KuaiRand (Kuaishou) Three released variants: KuaiRand-Pure is required, while KuaiRand-1k and KuaiRand-27k are bonus.	Short-video feed. 12 feedback signals (click / like / follow / comment / forward / long_view / play_time ...) plus a randomized-exposure intervention that supports counterfactual evaluation. Relevance label, task form and metrics are fixed by the organizers (pinned in the Starter Kit): the task treats <code>long_view</code> (native column) as the positive relevance label, ranks within each user's logged impressions (not full-catalog retrieval), and reports GAUC / nDCG@5 . Primary score = mean(GAUC, nDCG@5).	GAUC / nDCG@5	Pure: 1.4M interactions (27K users × 7.6K items). 1k: 11.7M. 27k: 322M.

Links: KuaiRand — <https://kuairand.com>

KuaiRand's randomized-exposure data also enables off-policy / counterfactual evaluation (OPE).

2.5 Deliverables

1. Written Project Description (via Devpost)

- Provide a clear written description of your project that includes:
 - How your solution addresses the problem statement
 - Development tools used (e.g. VSCode, Colab, Jupyter)
 - APIs used (e.g. OpenAI GPT-4o, Google Maps API)
 - Libraries and frameworks used (e.g. Hugging Face Transformers, PyTorch, scikit-learn, pandas)
 - Datasets and assets used (e.g. Google Local Reviews dataset, manually labelled data)

2. Public Code/GitHub Repository

- Submit a link to a public Code/GitHub repository containing:
 - Well-structured, commented code covering all components of your solution
 - A README file that includes:

- Project overview
- Setup and installation instructions
- Steps to reproduce your results
- A brief reflection on your solution's limitations and what you would improve given more time
- Team member contributions (if applicable, i.e. team participants, non-solo participants)

3. Run & Iteration Logs

- Submit the per-iteration log required in the Starter Kit (Run-log requirements), covering:
 - Hypothesis for that iteration — what the agent intended to try and why
 - The code diff applied
 - The resulting metrics (GAUC / nDCG@5 for the KuaiRand benchmarks)
 - Any error or recovery events encountered, and how the agent handled them
- A short summary reporting the number of manual interventions during the run (used to assess autonomy per Task Requirement 2)

4. Final Submission & Results Summary

- Submit your final model output/checkpoint for the required benchmark (KuaiRand-Pure), in the schema defined by the Starter Kit. If you also attempt the bonus benchmarks (KuaiRand-1k & KuaiRand-27k), submit their outputs as well for bonus scoring.
- A results table reporting your validation-best score for the required benchmark's metrics (KuaiRand-Pure GAUC / nDCG@5), and its absolute delta over the official baseline (per the Judging Criteria scoring formula); if you attempted the bonus benchmarks (KuaiRand-1k & KuaiRand-27k), include their GAUC / nDCG@5 results as well
- Reported resource usage required to reach the converged result: total token consumption (input + output) from the agent's LLM calls, the total **agent wall-clock** of the run, and the number of iterations used (out of the 50-iteration cap). Report GPU-hours as well if any GPU was used. These feed Feasibility & Practicality scoring.

2.6 Judging Criteria

Judging Criteria	Weight
Technical Execution	35%
Innovation & Problem Insight	20%

Impact & Relevance	20%
Feasibility & Practicality	15%
Presentation & Communication	10%
Final Event Only	

Technical Execution — Primary Metric & Robustness

Primary metric. We score the **converged result**, not the peak and not the intermediate trajectory. A run is considered converged when **validation score has not improved by more than $\epsilon = 0.002$ over the last $N = 3$ consecutive iterations**, *or* when the run hits the **50-iteration cap** or the **6 h wall-clock ceiling** — whichever comes first. The submission scored for ranking is the **validation-best checkpoint** at that point, evaluated **once on the hidden test set**. The agent develops only on train + validation; it never sees the hidden test set.

- **KuaiRand-Pure is the required benchmark** and determines 100% of the Primary metric score. **KuaiRand-1k and KuaiRand-27k are bonus benchmarks**: a strong result on either earns additional bonus points on top of the Primary metric score, but skipping them does not reduce the KuaiRand-Pure score.
- Per-dataset metrics: **KuaiRand-Pure / KuaiRand-1k / KuaiRand-27k $\rightarrow$ GAUC / nDCG@5**. Within each dataset, the score is the **equal-weighted average of each metric's *absolute* improvement over the official baseline** on the hidden test set. For every metric m :

代码块

```
1 delta(m) = score_agent(m) - score_baseline(m)
2 score_dataset = mean over m of delta(m)
```

- **Reading the numbers.** The metrics do not span $[0, 1]$. On the hidden test set, 27.1% of users have no positive label (their nDCG is 0 for any model) and 9.2% are all-positive, so a perfect ranking — using the true labels as the score — reaches only GAUC 1.0000 / nDCG@5 **0.7289** / primary **0.8645**. Random scoring sits at primary 0.4753. The official baseline's 0.5946 therefore already captures about 31% of the attainable range; judge progress against the 0.8645 ceiling, not against 1.0.

Robustness. Not judged by whether the agent ever hits a failure, but by **how it handles one** — recovering, retrying, or routing around a failed step (a code error, a timeout, an unexpected input) so that long iterative runs neither crash, stall, nor diverge before hitting the compute/wall-clock budget.

Innovation & Problem Insight

Judged on what the agent identified as worth trying and why — not on implementation.

- What the agent chose to target across the full algorithmic stack (features, model architecture, training strategy, evaluation loop, etc. — improvements are not limited to the model itself) and the reasoning behind that choice.
- Originality in drawing on published methods, papers, or public solutions — rewarding agents that go beyond naive baseline tweaks.

Impact & Relevance — Autonomy

Autonomy. How much of the improvement loop the agent drives on its own — proposing and testing changes based on its own evaluation of results, not just tuning the model architecture. Measured primarily by the **number of manual interventions** required to reach the converged result; fewer interventions score higher, with fully autonomous runs scoring highest. The fewer humans required, the more this reflects real acceleration of recommender-system R&D.

Feasibility & Practicality — Resource Consumption

How much it costs — in LLM usage and agent wall-clock — to reach the converged result. Two rules make this comparable: it is **scored only among submissions whose hidden-test primary score exceeds the official baseline**, and it is graded in **three coarse tiers** (low / medium / high consumption) rather than a continuous ranking. Without the quality gate the criterion would fight the Primary metric — an agent that stopped after three iterations would look cheapest and score worst.

- **Token consumption.** Total input + output tokens used by the agent's LLM calls across the run.
- **Agent wall-clock.** Total elapsed time of the agent run to reach the converged result. This replaces GPU-hours as the scored compute measure: on this benchmark the reference pipeline needs no GPU at all (about 28 min of single-core CPU for 100 iterations), so GPU-hours would be ~0 for most teams and would only penalise whoever happened to use a GPU. Report GPU-hours if any were used, but wall-clock is what is scored.

2.7 References

[1] J. S. Chan, N. Chowdhury, O. Jaffe, J. Aung, D. Sherburn, E. Mays, G. Starace, K. Liu, L. Maksin, T. Patwardhan, L. Weng, and A. Mądry, "MLE-bench: Evaluating Machine Learning Agents on Machine Learning Engineering," OpenAI, 2024. arXiv:2410.07095.

<https://doi.org/10.48550/arXiv.2410.07095>

[2] Z. Jiang, D. Schmidt, D. Srikanth, D. Xu, I. Kaplan, D. Jacenko, and Y. Wu, "AIDE: AI-Driven Exploration in the Space of Code," 2025. arXiv:2502.13138.

<https://doi.org/10.48550/arXiv.2502.13138>

[3] Y. Yamada, R. T. Lange, C. Lu, S. Hu, C. Lu, J. Foerster, J. Clune, and D. Ha, "The AI Scientist-v2: Workshop-Level Automated Scientific Discovery via Agentic Tree Search," 2025.

arXiv:2504.08066. <https://doi.org/10.48550/arXiv.2504.08066>

[4] H. Zhao, G. Cai, J. Zhu, Z. Dong, J. Xu, and J.-R. Wen, "Counteracting Duration Bias in Video Recommendation via Counterfactual Watch Time," KDD 2024. Code:

<https://github.com/hyz20/CWM> — *optional advanced reference, not the official baseline*. Its contribution is a censored-regression loss on watch time (a completed play means the true watch time was truncated by video length, so a one-sided loss is used instead of squared error). Note it ships no Recall implementation, reports nDCG@1/3/5 on a rebuilt `long_view2` label, and requires `torch==1.6.0`.

2.8 Appendix A. A Primer on Recommender Systems

This appendix gives participants without a recommender-systems background just enough to get started. It is a concept map plus an annotated reading list — not a textbook. Use it to understand the KuaiRand benchmarks and to know what to look up when you get stuck.

A.1 The Big Picture: The Recommendation Pipeline

A modern industrial recommender does not score every item directly. It runs a funnel of stages, each narrowing the candidate set:

代码块


1	Recall	→	Pre-ranking	→	Ranking	→	Re-ranking
2	millions		thousands		hundreds		final list

- **Recall / Retrieval:** cheaply retrieve a few thousand candidates from millions.
- **Pre-ranking:** a lightweight model trims the candidates further.
- **Ranking:** a heavy, accurate model scores each candidate. **This challenge mostly lives here.**
- **Reranking:** adjust the final ordering for diversity, business rules, and so on.

For this competition you mainly need the **ranking** stage. The KuaiRand benchmarks are ranking/prediction tasks, not full end-to-end pipelines.

The Recommendation Funnel

Millions of candidates in, a short final list out — every stage narrows the set.



A.2 Core Tasks: CTR and the Feedback Funnel

Most industrial ranking is framed as predicting the probability of user feedback:

- **CTR (Click-Through Rate)** — $P(\text{click} \mid \text{impression})$. The user saw the item; will they click?
- **CVR (Conversion Rate)** — $P(\text{conversion} \mid \text{click})$. The user clicked; will they convert (buy)? E-commerce background only; not a task in this challenge.
- **The funnel:** `impression` → `click` → `deeper engagement` (in e-commerce, → `conversion`). Because these stages are linked, two well-known problems arise:
 - **Sample selection bias:** the post-click signal is only observed on *clicked* items, yet must be predicted for *all* impressions.
 - **Data sparsity:** post-click signals such as `long_view` or `like` are far rarer than clicks.

KuaiRand has no purchase label, so CVR itself is never scored here. The funnel framing above is general background — note that in KuaiRand the scored label `long_view` is logged on *every* impression, not only on clicked ones, so classic sample selection bias does not apply directly to this challenge's task. Data sparsity still does, and the multi-feedback structure

(`click`, `like`, `follow`, `play_time` ...) makes ESMM-style multi-task modelling — see A.3 — a legitimate way to exploit the other signals as auxiliary tasks.

A.3 Multi-Task & Multi-Feedback Learning

Real users produce many signals (click, like, follow, comment, watch-time, and so on). Predicting them jointly — rather than training a separate model per signal — shares representations and tends to improve every task.

- Why it matters here: **KuaiRand** provides **12 feedback signals**, so a multi-task model can learn from several of them jointly even though only `long_view` is scored.
- The key idea is to balance *shared* parameters (which transfer useful knowledge across tasks) against *task-specific* parameters (which prevent conflicting tasks from hurting one another — the "seesaw" problem).

A.4 Evaluation Metrics

Metric	Intuition	Used for
AUC	Probability that a random positive is ranked above a random negative. Threshold-free and robust to class imbalance.	Scored in this challenge as GAUC — per-user AUC averaged with each user's positive count as the weight; users whose impressions are all-positive or all-negative are excluded.
NDCG	Quality of a <i>ranked list</i> , rewarding relevant items near the top (with a position discount).	Scored in this challenge as nDCG@5 . Users with no positive label score 0 and are included in the average.
Recall	Fraction of all relevant items that appear in the returned list.	Retrieval / coverage tasks — not scored here . Each user has only ~5 logged impressions in the evaluation split, so Recall@50 is 0.999+ for every model, including random scoring.

Offline vs. online: a higher offline metric does not always mean better real-world performance (because of distribution shift and feedback loops). This competition is evaluated offline, but it is worth knowing the gap exists.

A.5 Feature Engineering Basics

- **ID features:** user ID, item ID, category ID — high-cardinality discrete features.
- **Embedding:** map each discrete ID to a learnable dense vector. This is the foundation of all deep recommenders.

- **Feature crossing:** combine features (e.g. user $\times$ category) to capture interactions. Models such as FM and DeepFM automate this.

A.6 Annotated Reading List

[Hints: If you find reading the following material challenging or find you have missing backgrounds, you can use ChatGPT / Claude / ... to explain it to you.]

The goal here is only to understand **how a recommender system is structured** — the recall $\rightarrow$ ranking $\rightarrow$ re-ranking pipeline — and where the ranking stage (which this challenge targets) sits within it. You do **not** need to read a whole course; the introductory overview is enough. **Read just one of the following:**

- Google, *Recommendation Systems* (Machine Learning Crash Course), the **Overview** section — <https://developers.google.com/machine-learning/recommendation> A short, official overview of the pipeline. Note: Google calls the ranking stage "**scoring**" — this is the same thing as **ranking**, and it is the part this challenge focuses on.
- Wang Shusen, *Recommender Systems*, **Chapter 1 (Overview)** — <https://github.com/wangshusen/RecommenderSystem> The most beginner-friendly Chinese resource; the first chapter alone gives the full architecture.

3. Implement a GPU Kernel for a Transformer Layer



In response to some queries from our Early Bird participants, our engineers have provided updates to the problem statement to improve clarity and to support participants better.

Problem Statement last updated: 27 August 2026, 6:25PM

Added Appendix: Test Shapes

Updated  [torch_transformer_benchmark.py](#)



Technical Workshop Webinar with Q&A will be held on **28 Aug, 3:00 to 3:45pm.**

[Click here to join the webinar!](#)

3.1 Background

Transformer is a widely used neural network architecture in modern AI. It is the core structure behind many natural language processing, computer vision, speech, recommendation, and

large language model systems.

The main idea of Transformer is **self-attention**. Self-attention allows each token in a sequence to interact with other tokens directly. Compared with recurrent models, Transformer can process tokens in parallel, which makes it suitable for GPU acceleration.

Given an input sequence represented as a matrix:

$$X \in \mathbb{R}^{N \times d}$$

where N is the sequence length and d is the hidden dimension, the Transformer first projects the input into Query, Key, and Value matrices:

$$Q = XW_Q$$

$$K = XW_K$$

$$V = XW_V$$

The scaled dot-product attention is computed as:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

where d_k is the dimension of each attention head. The scaling factor $\sqrt{d_k}$ is used to prevent the dot-product values from becoming too large, which could make the softmax distribution unstable.

However, the computation of Transformer is expensive. Important operations include matrix multiplication, attention score calculation, softmax, normalization, and feed-forward layers. These operations may be limited by GPU compute throughput, memory bandwidth, cache efficiency, kernel launch overhead, and tensor core utilization.

In this competition, participants are asked to use AI-assisted methods to optimize the runtime efficiency of a Transformer structure on a given GPU model. The optimized implementation should improve performance while keeping the output numerically correct compared with the reference implementation.

Participants may consider optimization methods such as operator fusion, memory layout optimization, reduced-precision computation, tensor core usage, softmax optimization, and custom CUDA, Triton, TensorFlow, or PyTorch implementations.

The goal of this task is to explore how AI can help developers analyze Transformer workloads, identify bottlenecks, and generate more efficient implementations for specific GPU hardware.

3.2 Problem Statement

- Given a fixed formula of transformer layer, participants need to submit one or several GPU kernels that implement the layers that can pass the given test cases.

- The test cases would be written in pytorch or tensorflow and the participants can modify the layer implementation if they need, which means they can decide which parts of the layers should be fused into 1 kernel.
- The test case would compare the differences between the implementation of participants and the original pytorch/tensorflow implementation, the diff should be small enough (relative error < 0.02 , abs error < 0.002).
- The test cases would contain different shapes of input, including large/small batchsize, large/small sequence length, large/small dimensions, etc. The participants can choose different implementations for different shapes by adding shape checks in the implementation of layers. All the combinations of input shapes will be told to the participants.
- The use of AI tools is encouraged so that the participants can implement different kernels for different input shapes in limited time.
- Optimize & test your codes on your own machine. Different methods may be used to optimize the codes depending on the machine (GPU cards) you use.
- Provide a clear tech report including details on the AI skills/tools used to get bonus points.
- What participants need to do:
 - Download the benchmark scripts (choose either torch or tensorflow, one of them would be enough).
 - Implement the customized-implementation part and optimize it as fast as you can by AI or by hand.

```
class UserOptimizedTransformer(BaselineTransformer):
    """
    Replace this class with the optimized implementation.

    Requirements:
    1. Keep the forward signature unchanged.
    2. Return a tensor with shape [batch_size, seq_len, d_model].
    3. Keep compatible parameter names, or customize copy_model_weights().
    """

    def forward(
        self,
        x: torch.Tensor,
        valid_token_mask: Optional[torch.Tensor] = None,
    ) -> torch.Tensor:
        # ===== your codes here =====
        # Example optimization directions:
        # * torch.nn.functional.scaled_dot_product_attention
        # * torch.compile
        # * Triton/CUDA fused kernels
        # * fused LayerNorm / residual / FFN
        #
        # The default implementation calls the baseline so that this script
        # remains directly runnable before the optimized code is inserted.
        return super().forward(x, valid_token_mask)
        # =====
```

- Run the script on your own machine.
- Provide a clear tech report illustrating what the environment is (CPU, GPU, DISK, etc), what kind of optimizations you have done, and the final test results.

3.3 Constraints & Scope

Category	Constraints & Scope Details
In scope	AI-based code generation, GPU kernel fusion, profile tools usage, etc.
Out of scope	Production-ready deployment.

3.4 Available Resources / Data

You can download 1 of these, and run it on your own machine:

Torch Benchmark script

 [torch_transformer_benchmark.py](#)

Tensorflow Benchmark script

 [tensorflow_transformer_benchmark.py](#)

3.5 Deliverables

1. Written Project Description (via Devpost)

- Provide a clear written description of your project that includes:
 - How your solution addresses the problem statement
 - Development tools used (e.g. VSCode, Colab, Jupyter)
 - APIs used (e.g. OpenAI GPT-4o, Google Maps API)
 - Libraries and frameworks used (e.g. Hugging Face Transformers, PyTorch, scikit-learn, pandas)
 - Datasets and assets used (e.g. Google Local Reviews dataset, manually labelled data)

2. Public Code/GitHub Repository

- Submit a link to a public Code/GitHub repository containing:
 - Well-structured, commented code covering all components of your solution
 - A README file that includes:
 - Project overview
 - Setup and installation instructions
 - Steps to reproduce your results
 - A brief reflection on your solution's limitations and what you would improve given more time
 - Team member contributions (if applicable)

3. Demo Video

Submit a short video that:

- Demonstrates your solution working end-to-end (e.g. inference results, dashboard, model predictions)
- Is uploaded to YouTube and set to public visibility
- Is linked in your Devpost description
- Does not include third-party trademarks or copyrighted content without permission

Note for backend/NLP tracks: If a front-end interface is not applicable to your solution, a walkthrough video showing API usage, inference examples, or result analysis is accepted.

3.6 Judging Criteria

Judging Criteria	Definition	Weight
Technical Execution	The solution demonstrates strong engineering fundamentals, such as well-structured code, thoughtful architecture, and effective use of APIs or models. The demo runs reliably, and the technical complexity reflects deliberate, capable decision-making.	35%
Innovation & Problem Insight	The project demonstrates originality in both idea and approach. It stands out for the sharpness of its problem understanding — how clearly the team has framed the challenge, why it matters, and how directly the solution addresses it.	20%
Impact & Relevance	The project has clear potential to deliver value to real users or stakeholders — with meaningful reach, tangible benefit, and relevance that goes beyond solving for the hackathon prompt alone.	20%
Feasibility & Practicality	The solution is realistic and buildable beyond a prototype. The approach is technically and operationally sustainable — resource usage is proportionate, the architecture holds under real-world conditions, and the implementation is grounded rather than speculative.	15%
Presentation & Communication	[Final Event Only]: The team communicates their work with clarity. The pitch tells a coherent story; from problem to solution to potential, and the team is able to respond to questions with depth, demonstrating genuine understanding of their own project.	10%

3.7 Appendix

Test shapes:

#	Batch Size	QKV Dim	Heads	Seq Len	Layers	Causal	FFN Dim
1	64	128	4	128	4	TRUE	128
2	1	128	4	128	4	TRUE	128
3	4	128	4	128	4	TRUE	128
4	16	128	4	128	4	TRUE	128
5	128	128	4	128	4	TRUE	128
6	10000	128	4	128	4	TRUE	128
7	64	32	4	128	4	TRUE	32
8	64	1024	4	128	4	TRUE	1024
9	64	128	1	128	4	TRUE	128
10	64	128	2	128	4	TRUE	128
11	64	128	16	128	4	TRUE	128
12	64	128	4	32	4	TRUE	128
13	64	128	4	1024	4	TRUE	128
14	32	1024	16	100000	2	TRUE	1024

4. Shopping Copilot: AI Conversational Search and Recommendations

 **Technical Workshop Webinar with Q&A** will be held on **28 Aug, 4:00 to 4:45pm**.
[Click here to join the webinar!](#)

4.1 Background

Traditional e-commerce search engines heavily rely on static keyword matching, failing to capture the fluid shifts of genuine consumer psychology and the distinction between open-ended browsing and high-intent buying. In modern conversational commerce, constructing an intelligent agent that leverages dynamic context programming is critical to bridging the gap between ambiguous user queries and complex product catalogs. Solving this challenge directly impacts core industrial metrics.

4.2 Problem Statement

Participants are challenged to architect an intelligent, next-generation shopping agent capable of navigating real-world customer dynamics. Moving beyond rigid search filters, the engineered system must demonstrate deep cognitive understanding, runtime architectural agility, and commercial efficiency using the provided [Amazon dataset](#).

Specifically, the system should be built upon the following four core pillars:

I. Core Architecture: Intent Routing & Hybrid Pipeline

- **Dual-Track Routing:** Instantly detect the user's underlying intent—triggering a high-precision filter track for targeted "**Buying**" to lock hard constraints, and a diverse dense retrieval track for open-ended "**Browsing**" to unlock cross-category scenario matching.
- **Pipeline Base:** Construct an in-memory data stream featuring "**Multi-Route Retrieval** → **LLM Semantic Ranking**" (combining keyword, category, and vector similarity).

II. Dialog Strategy: Multi-Turn Scenario Evolution

- **Dynamic State Machine:** Build a robust conversational state tracker to gracefully handle dynamic **Information Accumulation** (incremental slots) and abrupt **Intent Override** (slot erasure and rewriting).
- **Proactive Guidance:** Trigger an immediate retrieval cutoff when facing **Over-Generality** (candidate pool overload) to actively generate structured, proactive clarification prompts that guide user convergence.

III. Self-Evolution: Dynamic Context Programming

- **Runtime Adaptation:** Leverage accumulated dialog history to perform **Personalized Context Distillation**, continuously updating short-term session states and long-term user profiles.
- **Adaptive Orchestration:** Utilize dynamic Context Programming to achieve runtime workflow re-orchestration and strategy alignment, ensuring the agent iteratively refines its own guidance logic.

IV. Evaluation Matrix: Product & Efficiency Metrics

Anchored on the final purchased record within the Amazon dataset, performance is quantified across three dimensions:

- **Coverage (Hit Rate@K):** Measures the catalog recall and boundary capability during the retrieval stage.
- **Precision (MRR / Top-K Hit Rate):** Evaluates the LLM's accuracy in pushing the exact purchased item to the absolute top of the recommendation list.

- **Efficiency (MTTC - Mean Turns to Conversion):** Heavy rewards systems that guide the user to the correct product in fewer interaction rounds, penalizing unnecessary conversational cognitive load.

4.3 Constraints & Scope

Category	Constraints & Scope Details
In scope	Designing highly sensitive intent-detection modules to split traffic into "Buying" and "Browsing" tracks. Implementing heterogeneous retrieval routing (weights, custom dynamic truncation, and slot decay over time). Engineering runtime-adaptive memory layers for personalized context distillation. Fine-tuning prompt strategies or local scoring logic for the LLM ranking stage to compress decision paths.
Out of scope	UI/UX Development (evaluated purely via automated backend APIs and headless pipelines). Training or full-parameter fine-tuning of base foundational LLMs. Deploying heavy external industrial vector DB clusters (must run entirely in-memory for light execution). Multi-Modal Processing (restricted strictly to text catalogs, structured metadata, and text dialogs).
Limits	Max Turns: Hard limit of 10 turns per session (forced termination and zero score if exceeded). Catalog Mutation: The Amazon product dataset is strictly read-only; no structural mutations or mock ASIN injections are allowed.
Allowed assumptions	Inputs are pre-cleaned text strings (no need to account for spelling correction, typos, or ASR noise). Product catalog, pricing, and category trees are static for the duration of the hackathon. Each session is simulated as an isolated single-user interaction (no multi-user concurrency stress needed).

4.4 Available Resources & Data

Participants receive a frozen and reproducible competition kit derived from the Amazon Reviews 2023 dataset.

Competition Data

- A frozen catalog containing 50,000 products from the Amazon Reviews 2023 `Clothing_Shoes_and_Jewelry` category.
- 200 labeled public development sessions for local testing and iteration.
- 800 additional sessions retained privately by the organizer for final evaluation.
- Public and private evaluation sessions use separate users and target products.

Participant Resources

- A weak BM25 starter Agent implemented in Python.
- A deterministic local evaluator for Hit Rate@10, MRR, MTTC, Efficiency, and the combined TechnicalScore.
- A published Python Agent interface and machine-readable API contract.
- Evaluation configuration, reproducible baseline results, data documentation, and submission rules.
- A SHA256 checksum file for verifying the downloaded catalog.

Participants can modify or replace the starter Agent while continuing to use the official local evaluator. The participant kit supports keyword retrieval, rule-based methods, dense retrieval, hybrid retrieval, reranking, local models, and external model APIs.

The organizer does not provide hosted model access, API keys, model tokens, or third-party API credits. A paid LLM is not required to complete the challenge. Teams that choose to use external services are responsible for their own credentials, usage limits, and costs, and must not publish secrets in their repositories.

Resources

- Participant repository:
<https://github.com/TechJam2026/techjam-conversational-search>
- Participant Kit Release:
<https://github.com/TechJam2026/techjam-conversational-search/releases/tag/participant-kit>
- Original data source and documentation:
<https://amazon-reviews-2023.github.io/>

The competition catalog and evaluation sessions are prepared and frozen by the organizer. Participants do not need to download or reconstruct the full upstream Amazon Reviews 2023 dataset.

4.5 Deliverables

1. Written Project Description (via Devpost)

- Provide a clear written description of your project that includes:
 - How your solution addresses the problem statement
 - Development tools used (e.g. VSCode, Colab, Jupyter)
 - APIs used (e.g. OpenAI GPT-4o, Google Maps API)
 - Libraries and frameworks used (e.g. Hugging Face Transformers, PyTorch, scikit-learn, pandas)
 - Datasets and assets used (e.g. Google Local Reviews dataset, manually labelled data)

2. Public Code/GitHub Repository

- Submit a link to a public Code/GitHub repository containing:
 - Well-structured, commented code covering all components of your solution
 - A README file that includes:
 - Project overview
 - Setup and installation instructions
 - Steps to reproduce your results
 - A brief reflection on your solution's limitations and what you would improve given more time
 - Team member contributions (if applicable, i.e. team participants, non-solo participants)

3. Demo Video

Submit a short video that:

- Demonstrates your solution working end-to-end (e.g. inference results, dashboard, model predictions)
- Is uploaded to YouTube and set to public visibility
- Is linked in your Devpost description
- Does not include third-party trademarks or copyrighted content without permission

Note for backend/NLP tracks: If a front-end interface is not applicable to your solution, a walkthrough video showing API usage, inference examples, or result analysis is accepted.

4.6 Judging Criteria

Judging Criteria	Definition	Weight
Technical Execution	The solution demonstrates strong engineering fundamentals, such as well-structured code, thoughtful	35%

	architecture, and effective use of APIs or models. The demo runs reliably, and the technical complexity reflects deliberate, capable decision-making.	
Innovation & Problem Insight	The project demonstrates originality in both idea and approach. It stands out for the sharpness of its problem understanding — how clearly the team has framed the challenge, why it matters, and how directly the solution addresses it.	20%
Impact & Relevance	The project has clear potential to deliver value to real users or stakeholders — with meaningful reach, tangible benefit, and relevance that goes beyond solving for the hackathon prompt alone.	20%
Feasibility & Practicality	The solution is realistic and buildable beyond a prototype. The approach is technically and operationally sustainable — resource usage is proportionate, the architecture holds under real-world conditions, and the implementation is grounded rather than speculative.	15%
Presentation & Communication	[Final Event Only]: The team communicates their work with clarity. The pitch tells a coherent story; from problem to solution to potential, and the team is able to respond to questions with depth, demonstrating genuine understanding of their own project.	10%

5. Robust Detection of AI-Generated Images Under Real-World Transformations

 **Technical Workshop Webinar with Q&A** will be held on **28 Aug, 5:00 to 5:45pm.**

Click here to join the webinar!

5.1 Background

Generative AI tools are making it easier than ever to create highly realistic synthetic images at scale. This creates new risks for online platforms, including misinformation, impersonation, fraud, and reduced trust in digital content. In practice, detection becomes even harder after

images are compressed, cropped, reposted, or lightly edited, so robust methods matter more than lab-only accuracy.

5.2 Problem Statement

We want participants to build a prototype that can distinguish AI-generated images from authentic images with strong robustness under realistic post-processing and redistribution scenarios. The goal is not only to achieve good detection performance on clean data, but also to maintain accuracy after transformations such as blur, compression, color adjustment, cropping, or rescaling. Solutions should present a clear technical approach, an evaluation strategy, and thoughtful discussion of trade-offs such as robustness, generalisation, and false positives.

Note: We consider robustness against a subset of the following augmentataions.

Transform	Parameters	Real-World Analog
JPEG Compression	quality = 90, 70, 50, 30	Social-media re-encode, messaging
Gaussian Blur	kernel σ = 0.5, 1.0, 2.0	Out-of-focus
Resize	scale $0.5\times$ / $0.25\times$ then upscale	Thumbnail generation
Gaussian Noise	σ = 0.02, 0.05, 0.10	Low-light sensor noise
Color Jitter	brightness/contrast/sat. $\pm 20\%$	Filter apps, auto-enhance
Center Crop	crop 80%	Profile-picture cropping, framing

5.3 Constraints & Scope

Category	Constraints & Scope Details
In scope	Image-level AIGC detection, robustness to common image transformations, feature engineering, model design, evaluation design, error analysis, and explainability ideas
Out of scope	Full production deployment, platform-wide moderation systems, and non-image modalities such as video or audio
Limits	Assume a hackathon-scale prototype, limited compute, and no access to internal production systems. Teams should optimise for a convincing proof of concept rather than a production-grade service. Note: Participants must use models with <2B parameters.
Allowed assumptions	

Teams may use public or properly licensed datasets, create their own transformed test cases, and make reasonable assumptions about deployment context as long as those assumptions are stated clearly.

5.4 Available Resources & Data

- Public or properly licensed image datasets for AIGC detection and image forensics.
- Self-created transformed samples using operations such as blur, compression, cropping, color adjustment, or rescaling.
- Public documentation for relevant machine learning and computer vision libraries.
- Datasets:
 - https://huggingface.co/datasets/saberztl/SID_Set
 - <https://www.kaggle.com/datasets/birdy654/cifake-real-and-ai-generated-synthetic-images>
 - <https://modelscope.cn/datasets/hy2628982280/WildFake/summary>
 - For this modelscope dataset, please translate it via the translation button before use:



Validation Dataset (for Demonstration Purposes Only):

We choose **a subset of WildFake** for participants to demonstrate their models' performance and track iterative improvements. This dataset serves only as a reference benchmark and will not contribute to the final score. **Do not use the following data during training.** Specifically:

Dataset		# Num
Non-AIGC	COCO val2017	4998
AIGC	DALL·E Advanced	8843

5.5 Expected Deliverables

1. Written Project Description (via Devpost)

- Provide a clear written description of your project that includes:
 - How your solution addresses the problem statement
 - Development tools used (e.g. VSCode, Colab, Jupyter)
 - Models or APIs used

- Libraries and frameworks used (e.g. Hugging Face Transformers, PyTorch, scikit-learn, pandas)
- Datasets and assets used

2. Public Code/GitHub Repository

- Submit a link to a public Code/GitHub repository containing:
 - Well-structured, commented code covering all components of your solution
 - A script that takes an image directory as input and outputs a confidence score for each image, indicating the likelihood that it is AIGC-generated. The output should be a JSON file containing `image_path` and `pred` for each image.
 - A README file that includes:
 - Project overview
 - Setup and installation instructions
 - Steps to reproduce your results
 - A brief reflection on your solution's limitations and what you would improve given more time
 - Team member contributions (if applicable, i.e. team participants, non-solo participants)

3. Demo Video

- Submit a short video that:
 - Demonstrates your solution working end-to-end (e.g. inference results, dashboard, model predictions)
 - Is uploaded to YouTube and set to public visibility
 - Is linked in your Devpost description
 - Does not include third-party trademarks or copyrighted content without permission

4. Robustness Evaluation Summary

- Include a compact table or visual summary comparing performance on clean images versus transformed images.

5. Error Analysis Note

- Highlight representative false positives, false negatives, and any trade-offs in the proposed approach.

5.6 Judging Criteria

--	--	--

Judging Criteria	Definition	Weight
Technical Execution	The solution demonstrates strong engineering fundamentals, such as well-structured code, thoughtful architecture, and effective use of APIs or models. The demo runs reliably, and the technical complexity reflects deliberate, capable decision-making.	35%
Innovation & Problem Insight	The project demonstrates originality in both idea and approach. It stands out for the sharpness of its problem understanding — how clearly the team has framed the challenge, why it matters, and how directly the solution addresses it.	20%
Impact & Relevance	The project has clear potential to deliver value to real users or stakeholders — with meaningful reach, tangible benefit, and relevance that goes beyond solving for the hackathon prompt alone.	20%
Feasibility & Practicality	The solution is realistic and buildable beyond a prototype. The approach is technically and operationally sustainable — resource usage is proportionate, the architecture holds under real-world conditions, and the implementation is grounded rather than speculative.	15%
Presentation & Communication	The team communicates their work with clarity. [Final Event Only]: The pitch tells a coherent story; from problem to solution to potential, and the team is able to respond to questions with depth, demonstrating genuine understanding of their own project.	10%

72-Hour Challenge & Project Submissions

Participants will have 72 hours to complete and submit their projects via [Devpost](#).

Refer to each track and problem statement to understand the specific deliverables required.

How To Submit Your Project

 **72-Hour Challenge & Submission Window: 29 August, 12pm to 1 September, 12pm**

All project submissions will be done via [Devpost](#), and can only be done during the **72-Hour Challenge** happening from **29 August, 12pm to 1 September, 12pm**. Here are some helpful resources from Devpost that could guide you:

- [How to enter a submission](#)
- [Know Your Submission Steps](#)
- [How to edit a submission](#)

When submitting, participants are reminded to ensure that all **deliverables** listed in the Problem Statement of your choice are included, and that all submission requirements have been addressed! 🍀

Judging Criteria

A panel of judges spanning across our product, engineering, security, operations, and design teams are charged with evaluating projects based on these key criteria:

Judging Criteria	Definition
Technical Execution	The solution demonstrates strong engineering fundamentals, such as well-structured code, thoughtful architecture, and effective use of APIs or models. The demo runs reliably, and the technical complexity reflects deliberate, capable decision-making.
Innovation & Problem Insight	The project demonstrates originality in both idea and approach. It stands out for the sharpness of its problem understanding — how clearly the team has framed the challenge, why it matters, and how directly the solution addresses it.
Impact & Relevance	The project has clear potential to deliver value to real users or stakeholders — with meaningful reach, tangible benefit, and relevance that goes beyond solving for the hackathon prompt alone.
Feasibility & Practicality	The solution is realistic and buildable beyond a prototype. The approach is technically and operationally sustainable — resource usage is proportionate, the architecture holds under real-world conditions, and the implementation is grounded rather than speculative.
Presentation & Communication	The team communicates their work with clarity. [Final Event Only]: The pitch tells a coherent story; from problem to solution to potential, and the team is able to respond to questions with depth, demonstrating genuine understanding of their own project.

Note: Tracks with different criteria will clearly specify expectations.

Prizes

Award	Prize
1st Place	SGD 15,000
2nd Place	SGD 8,000
3rd Place	SGD 5,000
4th Place	SGD 3,000
5th Place	SGD 3,000
People's Choice Award	SGD 500

People's Choice Award

Got a crowd-favorite project? The People's Choice Award puts the power in the hands of the community! After the submission period ends, public voting opens on [Devpost](#), and will run from **1 September, 3:00pm to 4 September, 3:00pm**. the project with the most verified votes takes home the prize. Here's how it works:

- Once your project is live, rally your crowd and share your Devpost submission link with your friends, peers, and network to gather votes!
- Anyone with a verified Devpost account can participate and vote (limited to one vote per person, per project).
- We strictly monitor voting. Using bots, fake accounts, or offering incentives/bribes for votes is prohibited and will result in immediate disqualification.
- In the rare event of a tie, the official judges' scores will determine the ultimate winner.

Build something awesome, share it with the world, and let the community decide!

Join Our Community



Be in the loop and get timely updates by joining our Telegram channel here:

@TikTokTechJam2026

FAQ

1. How do I register?

You must complete registration on **BOTH** the Devpost Page and Registration Form. You can register as a solo participant, or in a team with up to 5 members. Every member must complete both steps; i.e. submit the Registration Form and register on Devpost. Click the buttons below to register now:

Registration Form

Devpost

2. Do I really need to register via both Devpost and the Registration Form?

- Yes! Only those who register on both are eligible to participate.

3. Do all members in my team need to register via both Devpost and the Registration Form?

- Yes! Only those who register on both are eligible to participate.
- We require all members to register on both.

4. Why should I register early?

- There will be tips and sneak previews of the problem statements released via email weekly, and only those who register early will be able to receive these :)
- Early Bird entrants who register on **BOTH** the Devpost Page and Registration Form by **23 August 2026, Sunday, 11:59pm** will receive the Problem Statements earlier and enjoy a 48-hour head start.

5. Must I be a Singaporean citizen to join?

No. However, you must be a student at a university based in Singapore. For detailed information on eligibility, please refer to the [TikTok TechJam 2026 Rules](#).

6. Can I have a new team member joining my team after the registration deadline?

Yes. A new team member can officially join your team, as long as there are no more than 5 of you in the team. To do so, they must complete the [Registration Form](#), and must list you and all other members as teammates. The deadline for them to do so is **1 September 2026, 12pm**, which is also the Project Submissions Deadline on Devpost.

7. What if my team members change after registering?

Reach out to us at apac-earlycareers@tiktok.com, so we are able to take note of this.

8. Can I work on a project that is not part of any problem statement?

No. All teams will be required to build their solutions for the problem statements provided.

9. How do I submit my project?

All project submissions will be done via [Devpost](#)., and can only be done during the **72-Hour Challenge** happening from **29 August, 12pm to 1 September, 12pm**. Here are some helpful resources from Devpost that could guide you:

- [How to enter a submission](#)
- [Know Your Submission Steps](#)
- [How to edit a submission](#)

Be sure to ensure all **deliverables** listed in the Problem Statement of your choice are included, and that all **submission requirements** have been addressed! 🍀

10. Can I submit an existing project?

Your project must be either newly created, or, if your project existed prior to TikTok TechJam 2026, it must be significantly updated. Participants are reminded to address the specific problem statements, deliverables, and submission requirements.

11. Are late submissions allowed?

No. Please take note of the 72-Hour Challenge & Project Submissions Period happening from **29 August, 12pm to 1 September, 12pm**. Late submissions will not be considered.

12. If I am shortlisted as a Finalist, must I attend the Final in-person at the TikTok Singapore Office on 11 September 2026, 9am-6pm?

In general, yes, finalists are expected to participate in the in-person event. However, if you have questions and/or need additional support, please let us know via the Registration Form.



If you have other questions about this hackathon, reach out to us at apac-earlycareers@tiktok.com

Privacy Notice

This Campus Event Privacy Notice ("**Privacy Notice**") describes how TikTok Pte. Ltd, ("**we**" or "**us**") collects and processes your personal data in connection with your participation in the TikTok TechJam 2026 event (both virtual and onsite), "TikTok TechJam 2025" ("**Event**").

This Privacy Notice applies to the personal data of participants attending the Event, which may lead to potential employment opportunities. We think it is important that you understand how we use your personal data, and we take our obligations in this regard very seriously.

Controller of Personal Data

TikTok Pte. Ltd (a Singapore company) is the controller responsible for your personal data.

What categories of personal data will TikTok collect about me through this event?

TikTok may collect, process and use the following types of personal data about you when you have completed the registration form for the Event:

- **Identification data**, which includes your full name
- **Contact details**, which includes your email
- **Education and work experience**, which includes your current educational background, school, course/major
- **Other data**, which includes the information in your resume should you choose to submit your resume

Why does TikTok need my data?

We may collect and use your data to enable your participation in the Event and to engage you for potential employment opportunities.

If we use your personal data for recruitment processes, your personal data will be processed as explained under the [TikTok Global Applicant Privacy Notice](#).

Who might TikTok share my personal data with, and where?

As you may know, the TikTok entities organizing the Event are part of the global TikTok group, with offices located across the globe, which support event management, HR administration and provide services such as cloud storage. TikTok may also transfer your personal data to third parties outside of the TikTok group, including regulators, authorities, business partners, acquiring entities, service providers and other third parties, where necessary. Your personal data may be transferred to and processed in countries outside of Singapore where TikTok group and its service providers operate. We use appropriate safeguards to protect your data in all international transfers.

For more information on who we share your personal data with and where it is transferred to, you can contact us as set out below.

How long will TikTok retain my personal data?

It is TikTok's policy not to store personal data for longer than is necessary and to retain such data on a basis proportional to achieving the purposes prescribed above. This means that TikTok retains personal data for: (i) the period of time required for the purposes for which it was collected; (ii) any compatible and lawful purposes subsequently established; (iii) any new purposes to which you subsequently consent; and/or (iv) compliance with legal and regulatory requirements.

What rights do I have with respect to my personal data?

You may have specific rights under the data protection law. If you wish to learn more about these rights contact us as stated below.

These rights are not absolute and are subject to various conditions under applicable data protection and privacy legislation; and the laws and regulations to which TikTok is subject.

Who can I contact about this Notice?

If you have concerns or questions regarding this Notice or if you would like to exercise your rights as a data subject, you can contact us here: apac-earlycareers@tiktok.com